

# Lab 7a - CUDA base

<https://elly2023.smfi.unipr.it/mod/page/view.php?id=5580>

## OBIETTIVO

Sviluppare ed eseguire applicazioni GPU utilizzando CUDA-Toolkit.

## Teoria

### Architettura GPU vs CPU

La GPU è specializzata per il calcolo «data parallel» e i transistor del chip sono dedicati maggiormente al processamento dei dati piuttosto che al caching o al controllo di flusso.

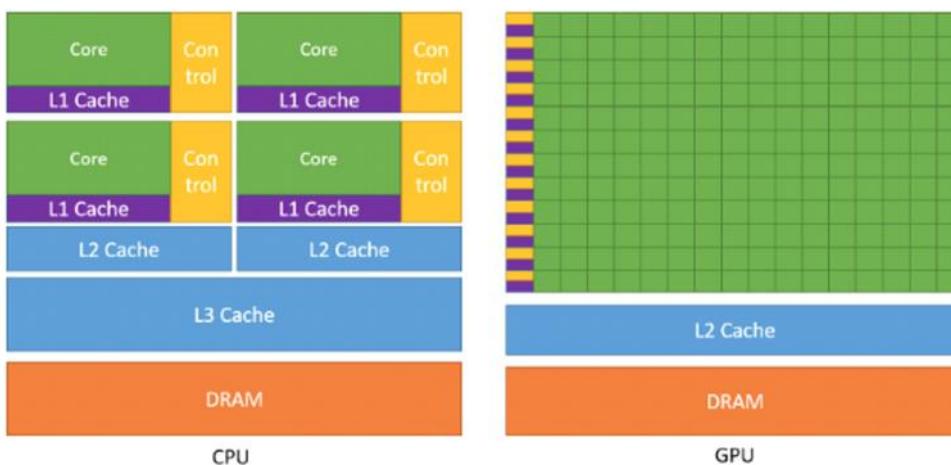


Figure 1: The GPU Devotes More Transistors to Data Processing

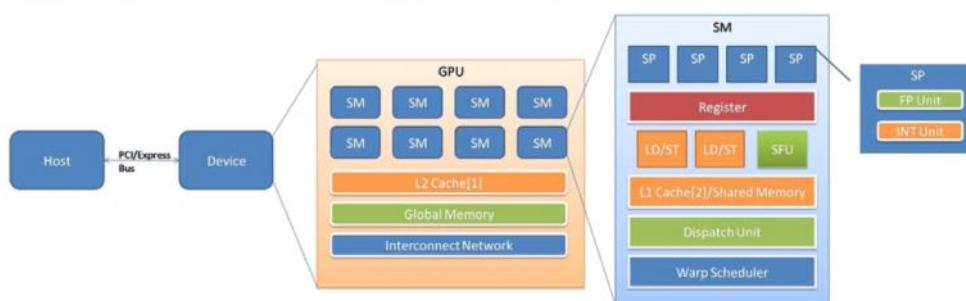
### Architettura GPU: SP, Warp e SM

L'unità elementare di calcolo è il **Thread**. Un thread viene eseguito da un Cuda core (**Streaming Processor – SP**).

I thread vengono eseguiti in parallelo in gruppi di 32 (Warp) su SP gestiti da uno **Steaming Multiprocessor (SM)**.

Ogni SM dispone di N SP (N=64 su P100, V100 e A100), un Warp scheduler, una memoria veloce (**Register**) privata per ogni thread, e una memoria veloce (L1 cache / **Shared memory**) condivisa tra i thread del SM.

Una GPU contiene un certo numero di SM (es. 56 su P100, 108 su A100) e un'ampia **memoria Globale** condivisa tra tutti i thread della GPU.



21/05/2024

5

## Esercizi Base

- Primo esercizio (dalla user guide del cluster HPC)

Ho copiato l'esercizio di esempio hello\_cuda.cu, dopodiché l'ho compilato ed eseguito.

In sintesi il programma modifica la stringa 'a' da "Hello" a "World!", sommando ad ogni carattere della prima un offset

per ottenere i caratteri della seconda. Stampando la stringa prima e dopo la modifica si ottiene in output "Hello World!"

### hello\_cuda.cu

```
/*
cp /hpc/share/samples/cuda/hello_cuda.cu .
module load cuda
nvcc hello_cuda.cu -o hello_cuda
*/

// This is the REAL "hello world" for CUDA!
// It takes the string "Hello ", prints it, then passes it to CUDA with an array
// of offsets. Then the offsets are added in parallel to produce the string "World!"
// By Ingemar Ragnemalm 2010

#include <stdio.h>

const int N = 7;
const int blocksize = 7;

// Kernel (memoria globale)
__global__
void hello(char *a, int *b)
{
    // threadIdx.x = indice thread nel blocco, va da 0 a 6
    a[threadIdx.x] += b[threadIdx.x];
}

int main()
{
    char a[N] = "Hello ";
    int b[N] = {15, 10, 6, 0, -11, 1, 0};

    // creo le strutture dati sull'host
    char *ad; // device a
    int *bd; // device b
    const int csize = N*sizeof(char);
    const int isize = N*sizeof(int);

    printf("%s", a);

    // creo le strutture dati sulle GPU
    cudaMalloc( (void**)&ad, csize );
    cudaMalloc( (void**)&bd, isize );
    cudaMemcpy( ad, a, csize, cudaMemcpyHostToDevice ); //copia 'ad' host-GPU (bloccante)
    cudaMemcpy( bd, b, isize, cudaMemcpyHostToDevice ); //copia 'bd' host-GPU (bloccante)

    // singola dimensione del blocco di thread utilizzato per eseguire il kernel (singolo blocco fatto di 7 thread)
    dim3 dimBlock( blocksize, 1 );

    // configurazione della griglia di blocchi (una sola cella (blocco), sia lungo x che y)
    dim3 dimGrid( 1, 1 );

    // chiamata del kernel (asincrona)
    hello<<<dimGrid, dimBlock>>>(ad, bd);

    // copia 'ad' GPU-host (bloccante)
    cudaMemcpy( a, ad, csize, cudaMemcpyDeviceToHost );

    // libera la memoria allocata sulla GPU
    cudaFree( ad );

    printf("%s\n", a);
    return EXIT_SUCCESS;
}
```

*Hello\_cuda.slurm*

```

#!/bin/sh
#< 1 node with 1 GPU
#SBATCH --partition=gpu
#SBATCH --qos=gpu
#SBATCH --nodes=1
#SBATCH --gres=gpu:p100:1
#SBATCH --time=0-00:30:00
#SBATCH --mem=4G
#< Charge resources to account
##SBATCH --account=<account>

module load cuda

nvcc hello_cuda.cu -o hello_cuda

echo "# SLURM_JOB_NODELIST : $SLURM_JOB_NODELIST"
echo "# CUDA_VISIBLE_DEVICES: $CUDA_VISIBLE_DEVICES"

./hello_cuda > hello_cuda.dat

```

*Output*

```

[martina.genovese@ui01 base]$ sbatch hello_cuda.slurm
Submitted batch job 2191799
[martina.genovese@ui01 base]$ cat hello_cuda.dat
Hello World!

```

- Altri esercizi

Ho copiato i file da "/hpc/home/roberto.alfieri/SHARE/cuda/base/", dopodiché ho testato tutti i programmi tramite lo script launch.slurm fornito

- [cuda\\_printf.cu](#) # in un kernel <<<10,10>>> ogni thread stampa le proprie coordinate

```

[martina.genovese@ui01 base]$ cat cuda_printf.dat
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0
Hello from block 0, thread 0
Hello from block 0, thread 1
Hello from block 0, thread 2
Hello from block 0, thread 3
Hello from block 0, thread 4
Hello from block 0, thread 5
Hello from block 0, thread 6
Hello from block 0, thread 7
Hello from block 0, thread 8
Hello from block 0, thread 9
Hello from block 6, thread 0
Hello from block 6, thread 1
Hello from block 6, thread 2
Hello from block 6, thread 3
Hello from block 6, thread 4
Hello from block 6, thread 5
Hello from block 6, thread 6
Hello from block 6, thread 7
Hello from block 6, thread 8
Hello from block 6, thread 9
Hello from block 3, thread 0
Hello from block 3, thread 1
Hello from block 3, thread 2
Hello from block 3, thread 3
Hello from block 3, thread 4
Hello from block 3, thread 5
Hello from block 3, thread 6
Hello from block 3, thread 7
Hello from block 3, thread 8
Hello from block 3, thread 9
Hello from block 9, thread 0
Hello from block 9, thread 1
Hello from block 9, thread 2
Hello from block 9, thread 3
Hello from block 9, thread 4
Hello from block 9, thread 5
Hello from block 9, thread 6
Hello from block 9, thread 7
Hello from block 9, thread 8

```

```
Hello from block 9, thread 9
Hello from block 5, thread 0
Hello from block 5, thread 1
Hello from block 5, thread 2
Hello from block 5, thread 3
Hello from block 5, thread 4
Hello from block 5, thread 5
Hello from block 5, thread 6
Hello from block 5, thread 7
Hello from block 5, thread 8
Hello from block 5, thread 9
Hello from block 1, thread 0
Hello from block 1, thread 1
Hello from block 1, thread 2
Hello from block 1, thread 3
Hello from block 1, thread 4
Hello from block 1, thread 5
Hello from block 1, thread 6
Hello from block 1, thread 7
Hello from block 1, thread 8
Hello from block 1, thread 9
Hello from block 7, thread 0
Hello from block 7, thread 1
Hello from block 7, thread 2
Hello from block 7, thread 3
Hello from block 7, thread 4
Hello from block 7, thread 5
Hello from block 7, thread 6
Hello from block 7, thread 7
Hello from block 7, thread 8
Hello from block 7, thread 9
Hello from block 4, thread 0
Hello from block 4, thread 1
Hello from block 4, thread 2
Hello from block 4, thread 3
Hello from block 4, thread 4
Hello from block 4, thread 5
Hello from block 4, thread 6
Hello from block 4, thread 7
Hello from block 4, thread 8
Hello from block 4, thread 9
Hello from block 2, thread 0
Hello from block 2, thread 1
Hello from block 2, thread 2
Hello from block 2, thread 3
Hello from block 2, thread 4
Hello from block 2, thread 5
Hello from block 2, thread 6
Hello from block 2, thread 7
Hello from block 2, thread 8
Hello from block 2, thread 9
Hello from block 8, thread 0
Hello from block 8, thread 1
Hello from block 8, thread 2
Hello from block 8, thread 3
Hello from block 8, thread 4
Hello from block 8, thread 5
Hello from block 8, thread 6
Hello from block 8, thread 7
Hello from block 8, thread 8
Hello from block 8, thread 9
```

- [sumArray.cu](#) # 2 vettori di float sull'host vengono copiati su GPU, sommati e copiato da GPU a host il \$  
[martina.genovese@u101 base]\$ cat sum\_array.dat  
#SLURM\_JOB\_NODELIST : wn41  
#CUDA\_VISIBLE\_DEVICES : 0  
Host> A: 1.0 2.0 3.0 4.0 5.0 6.0 7.0 8.0 9.0 10.0  
Host> B: 1.0 2.0 3.0 4.0 5.0 6.0 7.0 8.0 9.0 10.0  
sommaarray<<<DimGrid,DimBlock>>>  
Host> C: 2.0 4.0 6.0 8.0 10.0 12.0 14.0 16.0 18.0 20.0

- [matMul.cu](#) # 2 matrici NxN vengono copiate su GPU con 1 blocco 2D e moltiplicate

```
[martina.genovese@ui01 base]$ cat matMul.dat
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0

Host> A:
-1.0 0.0 1.0
2.0 3.0 4.0
5.0 6.0 7.0

Host> B:
1.0 2.0 3.0
4.0 5.0 6.0
7.0 8.0 9.0

matMul<<<DimGrid,DimBlock>>>

Host> C:
6.0 6.0 6.0
42.0 51.0 60.0
78.0 96.0 114.0
```

- [variabili.cu](#) # float A[16] viene copiato sulla GPU che somma elementi (riduzione somma) e torna il risultato

```
[martina.genovese@ui01 base]$ cat variabili.dat
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0
variabili.cu:1:1: error: unterminated comment
/*
^

Host> A:
-1.0 0.0 1.0
2.0 3.0 4.0
5.0 6.0 7.0

Host> B:
1.0 2.0 3.0
4.0 5.0 6.0
7.0 8.0 9.0

matMul<<<DimGrid,DimBlock>>>

Host> C:
6.0 6.0 6.0
42.0 51.0 60.0
78.0 96.0 114.0
```

- [managed.cu](#) # il vettore Vett e' managed; la GPU calcola il quadrato di ogni elemento

```
[martina.genovese@ui01 base]$ cat managed.dat
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0
0: Vett*Vett = 0
1: Vett*Vett = 1
2: Vett*Vett = 4
3: Vett*Vett = 9
4: Vett*Vett = 16
5: Vett*Vett = 25
6: Vett*Vett = 36
7: Vett*Vett = 49
8: Vett*Vett = 64
9: Vett*Vett = 81
10: Vett*Vett = 100
11: Vett*Vett = 121
12: Vett*Vett = 144
13: Vett*Vett = 169
14: Vett*Vett = 196
15: Vett*Vett = 225
16: Vett*Vett = 256
17: Vett*Vett = 289
18: Vett*Vett = 324
19: Vett*Vett = 361
20: Vett*Vett = 400
21: Vett*Vett = 441
22: Vett*Vett = 484
23: Vett*Vett = 529
24: Vett*Vett = 576
25: Vett*Vett = 625
26: Vett*Vett = 676
27: Vett*Vett = 729
28: Vett*Vett = 784
29: Vett*Vett = 841
30: Vett*Vett = 900
31: Vett*Vett = 961
```

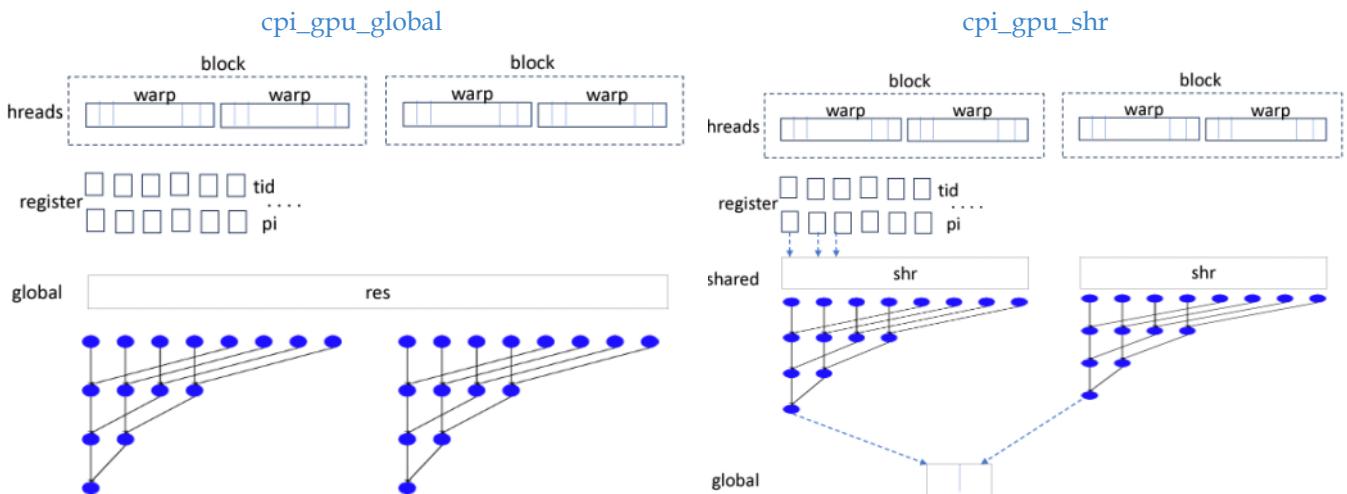
- race\_condition.cu # int a=0 viene copiato su GPU <<<100,100>>>, ogni thread somma 1 (race). Timing

```
[martina.genovese@ui01 base]$ cat race_condition.dat
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0
GPUtime elapsed: 0.000046 seconds
a=1
```

- sezione\_critica.cu # Sezione critica a livello warp o blocco con atomicCAS e atomicExch

```
[martina.genovese@ui01 base]$ cat sezione_critica.dat
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0
Block 0 thread 0 entering critical section
Block 1 thread 0 entering critical section
finito: 2
```

## Calcolo di Pigreco



Ho eseguito tramite cpi\_gp.slurm i programmi cpi\_gpu\_global e cpi\_gpu\_shr per il calcolo di pi-greco con GPU utilizzando la memoria global e la memoria shared.

## cpi\_gpu\_global.cu

```
#include <stdio.h>
#include <math.h>
#include <unistd.h> //optarg
#include <time.h>

#define BILLION 1000000000L;

void options(int argc, char * argv[]);

int n=500000000; // intervalli
int nblocks=128; // numero blocchi
int threadsPerBlock = 1024;
const double PI = 3.14159265358979323846264338327950288 ;

__global__ void add( float *res1 ) {

    long int tid = threadIdx.x + blockIdx.x * blockDim.x;

    double h = 1.0 / (double)(gridDim.x * blockDim.x);
    double x = h * ((double)tid - 0.5);
    double pi1 = (1.0 / (1.0 + x*x)); // f1
    res1[tid] = pi1 * 4 * h;

    __syncthreads();
    // for reductions, threadsPerBlock must be a power of 2 // because of the following code
    int i = blockDim.x/2;
    while ( i != 0 ) {
        if (threadIdx.x < i)
            res1[tid] += res1[tid + i];
        __syncthreads();
        i /= 2;
    }
}

/***********************/

int main(int argc, char **argv ) {

    options(argc, argv); /* optarg management */

    n=nblocks*threadsPerBlock;

    float* res1=(float*)malloc(n*sizeof(float));
    float *dev_res1;
    cudaMalloc( (void**)&dev_res1, n*sizeof(float) );

    struct timespec t1,t2,t3;
    double wtime, ktime;
    clock_gettime( CLOCK_REALTIME , &t1 );

    add<<<nblocks,threadsPerBlock>>>( dev_res1 );
    cudaDeviceSynchronize();
    clock_gettime( CLOCK_REALTIME , &t2 );

    cudaMemcpy( res1, dev_res1, n*sizeof(float), cudaMemcpyDeviceToHost );

    float total1=0;
    for (long int i=0;i<n;i+=threadsPerBlock) total1+=res1[i];

    clock_gettime( CLOCK_REALTIME , &t3 );

    wtime = (double) ( t3.tv_sec - t1.tv_sec )
        + (double) ( t3.tv_nsec - t1.tv_nsec )
        / BILLION;

    ktime = (double) ( t2.tv_sec - t1.tv_sec )
        + (double) ( t2.tv_nsec - t1.tv_nsec )
        / BILLION;

    fprintf(stderr,"#intervals blocks pi error wtime(s) ktime(s) \n");
    fprintf(stderr,"CUDA, %d, %d, %d, %.10f, %.10e, %.4f, %.4f \n",
        n, nblocks, threadsPerBlock, total1, fabs(total1 - PI), wtime, ktime);
    cudaFree( dev_res1 );
}
```

```

        return 0;
    }

/***********************/

void options(int argc, char * argv[])
{
    int i;
    while ( (i = getopt(argc, argv, "t:b:h")) != -1) {
        switch (i)
        {
        case 'b': nblocks      = strtol(optarg, NULL, 10); break;
        case 't': threadsPerBlock = strtol(optarg, NULL, 10); break;
        case 'h': printf ("\n%s [-b blocks] [-h]",argv[0]); exit(1);
        default:   printf ("\n%s [-b blocks] [-h]",argv[0]); exit(1);
        }
    }
}

```

## cpi\_gpu\_shr.cu

```

#include <stdio.h>
#include <math.h>
#include <unistd.h> //optarg
#include <time.h>

#define BILLION 1000000000L;

void options(int argc, char * argv[]) ;

int n=500000000; // intervalli
int nblocks=128; // numero blocchi
int threadsPerBlock = 1024;
const int max_threadsPerBlock = 1024;
const double PI = 3.14159265358979323846264338327950288 ;

__global__ void add( float *res1 ) {
    __shared__ float shr1[max_threadsPerBlock];

    int tid = threadIdx.x + blockIdx.x * blockDim.x;

    double h = 1.0 / (double)(gridDim.x * blockDim.x);
    double x = h * ((double)tid - 0.5);
    double pi1 = (1.0 / (1.0 + x*x)); // f1
//    double pi2 = sqrt(1-x*x); // f2

    shr1[threadIdx.x] = pi1 * 4 * h;

    __syncthreads();
    // for reductions, threadsPerBlock must be a power of 2 // because of the following code
    int i = blockDim.x/2;
    while ( i != 0) {
        if (threadIdx.x < i)
            shr1[threadIdx.x] += shr1[threadIdx.x + i];
        __syncthreads();
        i /= 2;
    }

    if (threadIdx.x==0)      res1[blockIdx.x] = shr1[threadIdx.x];
}

/***********************/

int main(int argc, char **argv ) {

    options(argc, argv); /* optarg management */

    n=nblocks*threadsPerBlock;

    float* res1=(float*)malloc(nblocks*sizeof(float));
    float *dev_res1;
    cudaMalloc( (void**)&dev_res1, nblocks*sizeof(float) );

    struct timespec t1,t2,t3;

```

```

double wtime,ktime;
clock_gettime( CLOCK_REALTIME , &t1 );

add<<<nblocks,threadsPerBlock>>>( dev_res1);
cudaDeviceSynchronize();

clock_gettime( CLOCK_REALTIME , &t2 );

cudaMemcpy( res1, dev_res1, nblocks*sizeof(float), cudaMemcpyDeviceToHost );

float total1=0;
for (int i=0;i<nblocks;i++)    total1+=res1[i];

clock_gettime( CLOCK_REALTIME , &t3 );

wtime = (double) ( t3.tv_sec - t1.tv_sec )
+ (double) ( t3.tv_nsec - t1.tv_nsec )
/ BILLION;

ktime = (double) ( t2.tv_sec - t1.tv_sec )
+ (double) ( t2.tv_nsec - t1.tv_nsec )
/ BILLION;

fprintf(stderr,"#intervals blocks pi error wtime(s) ktime(s)\n");
fprintf(stderr,"CUDA, %d, %d, %d, %.10f, %.10e, %.4f, %.4f \n", n, nblocks, threadsPerBlock, total1, fabs(total1 - PI),
wtime, ktime);
cudaFree( dev_res1 );

return 0;
}

/********************************************/


void options(int argc, char * argv[])
{
int i;
while ( (i = getopt(argc, argv, "t:b:h")) != -1) {
switch (i)
{
case 'b': nblocks = strtol(optarg, NULL, 10); break;
case 't': threadsPerBlock = strtol(optarg, NULL, 10); break;
case 'h': printf ("\n%s [-b blocks] [-h]",argv[0]); exit(1);
default:   printf ("\n%s [-b blocks] [-h]",argv[0]); exit(1);
}
}
}

```

### File di output

```

[martina.genovese@ui01 cpi]$ more cpi_gpu.slurm.o2191772
#SLURM_JOB_NODELIST : wn41
#CUDA_VISIBLE_DEVICES : 0
#intervals blocks pi error wtime(s) ktime(s)
CUDA, 1048576, 1024, 1024, 3.1415929794, 3.2584135923e-07, 0.0026, 0.0001
#intervals blocks pi error wtime(s) ktime(s)
CUDA, 1048576, 1024, 1024, 3.1415929794, 3.2584135923e-07, 0.0002, 0.0002

```

Ho poi creato lo script cpi\_gpu\_scaling.slurm in modo da valutare la scalabilità dei parametri significativi.

### cpi\_gpu\_scaling.slurm

```

#!/bin/sh

#SBATCH --output=%x.o%
#SBATCH --partition=gpu
#SBATCH --qos=gpu
#SBATCH --gres=gpu:p100:1
#SBATCH --time=0:00:05:00
#SBATCH --mem=12G

#stampa il nome del nodo assegnato e argomenti
echo "#SLURM_JOB_NODELIST : $SLURM_JOB_NODELIST"
echo "#CUDA_VISIBLE_DEVICES : $CUDA_VISIBLE_DEVICES"

module load cuda

threadsPerBlock=1024
# nblocks=1024

```

```

nvcc cpi_gpu_global.cu -o cpi_gpu_global

for nblocks in 256 512 1024 2048 4096 8192
do
    cpi_gpu_global -b $nblocks -t $threadsPerBlock 2>> cpi_gpu_scaling_global.dat
done

nvcc cpi_gpu_shr.cu -o cpi_gpu_shr

for nblocks in 256 512 1024 2048 4096 8192
do
    cpi_gpu_shr -b $nblocks -t $threadsPerBlock 2>> cpi_gpu_scaling_shr.dat
done

```

Ho infine creato il programma in python cpi\_gpu\_scaling.py per confrontare le prestazioni delle varie implementazioni del programma.

### cpi\_gpu\_scaling.py

```

#!/usr/bin/env python2

import matplotlib
matplotlib.use('Agg') # Backend per PNG
import matplotlib.pyplot as plt
import pandas as pd

dfG = pd.read_csv("cpi_gpu_scaling_global.dat", names=["mem", "intervals", "nblock", "threadxblock", "pi", "error", "wtime",
"ktime"])
dfS = pd.read_csv("cpi_gpu_scaling_shr.dat", names=["mem", "intervals", "nblock", "threadxblock", "pi", "error", "wtime", "ktime"])

print(dfG)
print(dfS)

plt.subplot(2,1,1)

plt.title('CPI GPU Scaling')
plt.grid()
plt.xlabel('blocks')
plt.ylabel('time')
plt.yscale('log')
plt.plot(dfG.nblock, dfG.wtime, '-o', label="GPU Global")
plt.plot(dfS.nblock, dfS.wtime, '-o', label="GPU Shared")
plt.legend(shadow=True, loc="best")

plt.subplot(2,1,2)

plt.grid()
plt.xlabel('blocks')
plt.ylabel('speedup')
plt.plot(dfG.nblock, dfG.wtime.iloc[0] / dfG.wtime, '-o', label="GPU Global")
plt.plot(dfS.nblock, dfS.wtime.iloc[0] / dfS.wtime, '-o', label="GPU Shared")
plt.plot(range(1,60),range(1,60),'-r', label='Ideal')

plt.legend(shadow=True, loc="best")

plt.savefig('cpi_gpu_scaling.png')

```

*cpi\_gpu\_scaling.png*

...